


INTRO TO DATA SCIENCE

LECTURE 5: LINEAR MODELS & LINEAR REGRESSION

WHERE ARE WE ON THE DATA SCIENCE ROAD-MAP?



Finding the minimum of the cost function determines values for θ that optimally (in the sum of squared errors sense) model the training set

A second algorithm for finding the minimum of the cost function are via the Normal Equations

Standard minimization problem

Take the derivative and set equal to zero

Using the training set can derive a group of simultaneous equations

This is not an iterative algorithm, rather it is purely analytical and solves for the final values of theta directly

The matrix containing the training data is called 'The Design Matrix'

Although there is no strict need to feature scale, by force of habit I generally always scale my features

$$\vec{\theta} = (\vec{X}^T \vec{X}) \vec{X}^T \vec{y}$$

where $\vec{\theta}$ is the vector of model parameters, \vec{X} is an augmented matrix containing the training data (The Design Matrix). It is augmented in the sense that it contains the intercept parameter. y are the outputs of the training data

KEY CONCEPTS - LINEAR MODELS AND MULTIVARIATE LINEAR REGRESSION

$$y = \theta_0 + \theta_1 x$$

Univariate Linear Regression

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 \dots \theta_n x_n$$

Multivariate Linear Regression

Extend the number of features beyond 1.

KEY CONCEPTS - ASSUMPTIONS FOR LINEAR REGRESSION

1. Linear relationship - the relationship between features and output is linear
2. Multivariate normality - the features are identically distributed
3. No or little multicollinearity - the features are independent
4. No auto-correlation - the residuals are independent of each other
5. Homoscedasticity - the variance in the residuals is 'constant'

If the features are not independent and exhibit multi-collinearity, the normal equations can become ill-posed

Adding a regularizer solves this problem

1. Gradient Descent

- i. An iterative algorithm
- ii. Need to know the cost function
- iii. Start with an initial random guess for the model parameters
- iv. At each iteration we alter the parameters a little to reduce the cost function
- v. Stop when a minimum of the cost function is reached
- vi. The training data must be scaled before we used this algorithm
- vii. We need to determine a hyper-parameter, called the learning rate
- viii. Why do you need to know about this algorithm??

2. Normal Equations

- i. It is possible to solve for the model parameters directly!!
- ii. Enter the training data (inputs and outputs) into a set of equations and determine the values of theta
- iii. No need to scale the inputs
- iv. No hyper-parameter to estimate
- v. Involves the calculation of the inverse of a matrix
- vi. As the dataset gets large this becomes very computationally expensive
- vii. Equations can become unstable if linear regression assumptions are violated